

Speech time-scale modification using GANs

By Jayin Khanna

Supervisor: Dr Prasanta Kumar Ghosh

Mentors: Haritha Neelapuja & Jesuraj

1. Defining and the Problem
2. ScalerGAN
3. Prerequisite: GANs
4. Architecture

Problem statement

Alter speech rate (slow down/speed up) while preserving:

- Speaker identity i.e.
 - Timbre
 - Pitch
- Intelligibility and,
- Naturalness.

Upshot: **scaling the time axis** of the signal while attempting to keep the **spectral content** (pitch, timbre) intact.

Challenges:

1. Traditional signal processing methods (e.g. PSOLA, WSOLA, Phase Vocoder) cause artifacts at extreme rates.
2. No supervised data: Parallel utterances at different rates are unavailable.

Goal: Develop a deep learning solution for high-quality TSM without paired training data.

Proposed solution: ScalerGAN

Unsupervised GAN framework inspired by CycleGAN and InGAN, adapted for speech spectrograms:

A new deep learning algorithm for TSM that can speed up or slow down speech signals at a given rate.

Key Components:

1. **Generator (G):**
 - Input: Speech spectrogram + desired rate r .
 - Output: Time-scaled spectrogram.
2. **Multi-Scale Discriminator (D):**
 - 5 sub-discriminators at different spectrogram scales.
 - Classifies "real vs. synthetic" spectrograms.
3. **Decoder:**
 - Reconstructs original spectrogram from time-scaled output using G with rate r^{-1} (cycle consistency).

Generative Adversarial Networks (GANs): Core Idea

The Min-Max Game

A GAN consists of two neural networks locked in a competitive game:

Real: 1
Fake: 0

- **Generator (G):** The forger. Tries to create fake data that looks real.
- **Discriminator (D):** The detective. Tries to distinguish real data from fakes.

Their goals are directly opposed, formalized by this Min-Max Objective:

The value function is:

$$V(D, G) = \mathbb{E}_{x \sim p_{data}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))]$$

D(x): D's estimate that real data x is real (goal: ~1).

D(G(z)): D's estimate that fake data G(z) is real (goal: ~0 for D, ~1 for G).

The Ultimate Goal: Distribution Matching

The generator doesn't just memorize data; it learns to model the underlying probability distribution (p_{data}) of the real data.

- **Initial State:** Generator outputs random noise from a simple distribution (e.g., Gaussian).
- **Final State:** Generator transforms noise into data that lies on the same space as the real data.
- At convergence: Generator's distribution $p_g \approx p_{data}$, and the discriminator is forced to guess ($D(x) = 0.5$ everywhere).

Training of GANs

Step-by-Step Training Loop

For each training iteration:

1. **Sample Data:**
 - o Get a batch of real data x and a batch of random noise vectors z .
2. **Generate Fakes:**
 - o Forward pass through Generator: $G(z) = x_{\text{fake}}$.
3. **Update the Discriminator (D):**
 - o Zero D's gradients.
 - o **Forward Pass:** Compute $D(x)$ (real) and $D(x_{\text{fake}})$ (fake).

Calculate G's Loss:

$$L_G = -\log(D(x_{\text{fake}}))$$

- o This loss is high if D correctly identifies the fake ($D(x_{\text{fake}})$ is low).
- o Backpropagate L_G to calculate gradients for G's parameters.
- o Update G's weights to become a better forger.

Calculate D's Loss:

$$L_D = -[\log(D(x)) + \log(1 - D(x_{\text{fake}}))]$$

- o This loss is high if D is wrong (e.g., if $D(x_{\text{fake}})$ is high).
- o Backpropagate L_D to calculate gradients for D's parameters.
- o Update D's weights to become a better detective.

4. **Update the Generator (G):**
 - o Zero G's gradients.
 - o Generate new fakes: $G(z) = x_{\text{fake}}$ (Forward pass again).
 - o Re-run D on new fakes: $D(x_{\text{fake}})$ (D is now frozen).

Generative adversarial Networks

Framework: Two neural networks contesting:

1. **Generator (G):** Creates synthetic data (e.g., spectrograms).
2. **Discriminator (D):** Distinguishes real vs. synthetic data.

Training Dynamics:

1. **Generator (G):** Creates synthetic data from random noise.
2. **Discriminator (D):** Distinguishes real data from synthetic data.

GAN Objective Function:


$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

1. Discriminator maximises the probability of assigning the correct labels
2. Generator minimizes the log probability that the discriminator labels generated image as fake

GAN training


Repeat for each training iteration:

1. Sample minibatch of real data:
 $x_1, \dots, x_n \sim p(\text{data})$
2. sample noise vectors:
 $z_1, \dots, z_n \sim p_z$
3. Compute $D(x_i)$ and $D(G(z_i))$
4. Update Discriminator using gradient ascent:
 $\nabla D[\log D(x) + \log(1 - D(G(z)))]$
5. Update Generator using gradient descent:
 $\nabla G[-\log D(G(z))]$



* To train the discriminator:

→ keep Θ constant, $D = \{x_1, x_2, \dots, x_n\}$




$$B = \{x_1, \dots, x_{B_1}\} \subset D$$

One update: $\omega^{t+1} \leftarrow \omega^t + \alpha \nabla_\omega J_{GAN}$

$$J_{GAN}(\Theta, \omega) = \left[\frac{1}{B_1} \sum_{i=1}^{B_1} \log D_w(x_i) + \frac{1}{B_2} \sum_{j=1}^{B_2} \log(1 - D_w(g_\theta(z_j))) \right]$$

* To train the Generator Network:



Update θ only
with w Constant

$$J_{GAN} = \left[\frac{1}{B_2} \sum_{j=1}^{B_2} \log(1 - D_w(g_\theta(z_j))) \right]$$

$$\theta^{t+1} \leftarrow \theta^t - \alpha_2 \nabla_\theta J_{GAN}(\theta, w)$$

Figure 3: Train Generator network from notes

ScalerGAN

Goal: Alter speech speed (slow down/speed up) while preserving:

- Speaker identity (timbre, pitch).
- Intelligibility and naturalness

Input:

A Mel-spectrogram **matrix $X \in \mathbb{R}^{(F \times N)}$**

- F : number of frequency bins (e.g., 80 Mels)
 - N : number of time frames
- Desired scaling rate r ($r > 1$: speed up, $r < 1$: slow down)

Output:

A scaled Mel-spectrogram **Y $\in \mathbb{R}^{(F \times M)}$**

- $M = \lfloor N \cdot r \rfloor$ new time frames

Learn a generative function $G : (X, R) \rightarrow X^*$ that given a finite-length sequence and the desired rate r will generate a finite sequence with a scaled duration according to the given rate r .

The most trivial implementation of the function G is by resampling the original signal x
Problems? Loses speaker identity Distorts pitch Produces artifacts

Why Mel-Spectrograms for Speech TSM

What is a Mel Spectrogram?

A visual representation of sound that shows:

- **Time** on the horizontal axis (when sounds occur).
- **Frequency** on the vertical axis (which pitches are present).
- **Intensity/Energy** shown by color (brighter = louder).



How is it Made?

The process converts a raw waveform into a Mel Spectrogram:

1. **Framing & Windowing:** Chop the audio into short, overlapping segments (e.g., 25ms).
2. **Fourier Transform (STFT):** Calculate the frequency spectrum for each segment.
3. **Mel Scaling:** Apply a filter bank that warps the frequencies to the **Mel scale**, which mimics human hearing.
4. **Log Compression:** Convert amplitude to a logarithmic (dB) scale, as we perceive loudness logarithmically.

Why Use Mel-Spectrograms as inputs?

5. **Perceptually Relevant:** The Mel scale emphasizes crucial speech frequencies (like formants) and de-emphasizes less important ones, just like the human ear does.
6. **Reduced Dimensionality:** They provide a compact, information-rich summary of the audio, making the model's job more efficient than using raw waveforms or linear spectrograms.



Overall Architecture of the Model

1. Generator (G)

Input: A speech spectrogram \bar{x} and a desired time-scaling rate r

Operation:

- Applies a linear interpolation on the time axis by factor r
- Passes the interpolated spectrogram through a U-Net to refine and enhance it

Output: A time-scaled spectrogram \bar{y} corresponding to the new rate r

2. Multi-Scale Discriminator (D)

Goal: Distinguishes between real and generated (time-scaled) spectrograms

Design:

1. Consists of multiple sub-discriminators, each operating at a different temporal scale
2. Uses downsampling and patch-level classification to assess local realism at various resolutions

Training Signal: Guides the generator to produce spectrograms that are indistinguishable from real ones

3. The Cycle-Consistency Decoder

Input: Takes the time-scaled spectrogram \bar{y} back and the inverse of the time rate i.e. r^{-1}

Process:

- The time-scaled spectrogram \bar{y} is passed back to the same generator G with an inverse rate $1/r$
- Generates a reconstructed spectrogram \bar{x}'

Objective: The reconstructed spectrogram should match the original input:

$$G(G(X, r), 1/r) \approx X$$




Fig. 1. The ScalerGAN model. The generator G gets a spectrogram \bar{x} and a rate r as input, and generates a new time-adjusted spectrogram \bar{y} . During training, the resulting spectrogram is further processed by a discriminator D that gets as input either \bar{x} or \bar{y} and decides whether the input is real or fake. Parallel to that, a decoder gets \bar{y} as input and tries to reconstruct \bar{x} back from it, where the decoder is implemented using the same generator G . At inference time, we use a vocoder to convert the spectrogram into a waveform.

The Generator

What it does

- Core transformation unit
- Takes an input spectrogram and target rate
- Produces a **time-scaled spectrogram $Y = G(X, r)$**

How it works

- **Input:** Mel-spectrogram X and desired rate r
- **Architecture:** U-Net with skip connections
- **Output:** New spectrogram Y , length scaled by r ($M = \lfloor N \cdot r \rfloor$)
- **Key Function:** Reshapes the temporal axis while preserving frequency content & harmonic structure

U-Net Architecture

1. Encoder: Hierarchical Feature Extraction

- **Purpose:** Projects input spectrogram into a latent feature space
- **Process:** Strided convolutions \downarrow temporal resolution, \uparrow channel depth (feature richness)
- **Why for TSM?** Learns abstract, contextual representations (phonemes, syllables)

2. Bottleneck: High-Level Representation Learning

- **Purpose:** Core transformation zone for reorganizing features
- **Process:** 6 residual blocks apply non-linear transformations
- **Why for TSM?** Interprets scaling rate r and plans temporal expansion/contraction
- Residual connections \rightarrow stable deep training

3. Decoder: Temporal Upsampling & Reconstruction

- **Purpose:** Synthesizes output spectrogram of target length M
- **Process:** Transposed convolutions \uparrow temporal resolution, \downarrow channel depth
- **Why for TSM?** Performs temporal resampling while reconstructing structure from bottleneck features

4. Skip Connections: Preserving Spectral Fidelity

- **Purpose:** Bypass bottleneck \rightarrow transfer encoder features directly to decoder
- **Why critical for TSM?**
 - Without: outputs are temporally correct but spectrally blurry
 - With: outputs retain sharp harmonic structures & formants \rightarrow naturalness + speaker identity preserved

The Discriminator

Core Concept: Why Multi-Scale?

- **Goal:** Ensure generated spectrogram $Y = G(X, r)$ is realistic across all time–frequency resolutions
- **High-Resolution (original scale):**
 - Judges fine details → formant sharpness, consonant clarity, harmonic texture
 - Catches local blurring/noise
- **Low-Resolution (downscaled):**
 - Judges macro-structure → rhythm, prosody, energy contours
 - Catches unnatural pauses, rhythm issues
- **Inspiration:** From InGAN → realism across spatial frequencies; directly applicable since spectrograms are 2D time–frequency images

Architecture: How It's Implemented

- **Discriminator D** = ensemble of K sub-discriminators $\{D_1, D_2, \dots, D_K\}$
- **Input Processing:**
 - $D_1 \rightarrow$ original resolution
 - $D_2 \dots D_K \rightarrow$ progressively downsampled (avg pooling, factors 2, 4, ...)
- **Sub-Discriminator Design:**
 - Each D_K is PatchGAN-style classifier
 - Strided convolutions for downsampling
 - Outputs 2D matrix of probabilities (not a single scalar)

PatchGAN. Why?

- Judges spectrograms patch-by-patch, not globally
- Each matrix element = probability that local patch is real
- Enforces local + global coherence → suppresses artifacts in structured data like spectrograms

SPOT THE 3 DIFFERENCES



1.200 × 675

Multi-Scale Discriminator: Exact Data Flow & Patchwise Analysis

Initial Conditions:

- Real Spectrogram (Target): X_{real} ,
- of shape [1, 80, 256] (Channels, Freq, Time).
- Fake Spectrogram (Input): $Y_{\text{fake}} = G(X, 1.5)$ of shape [1, 80, 384].
- Scaling Factors: $\ln \{1, 1.2, 1.2^2, 1.44, 1.2^3, 1.2^4\}$.

These are the factors for spatial downscaling.

Step 1: Define Target Size: The target size for the input to discriminator D_s is the **original spatial size divided by s** .

- For X_{real} (256×80): $T_s(X_{\text{real}})$ resizes it to [256/ s , 80/ s].
- For Y_{fake} (384×80): $T_s(Y_{\text{fake}})$ resizes it to [384/ s , 80/ s].
- Example: For $s=1.44$, X_{real} is resized to ~[178, 56] and Y_{fake} to ~[267, 56].*

The Transformation Operator (T): This is typically a **bilinear interpolation** (or bicubic). Its purpose is to **smoothly and differentially downscale** the spectrogram, simulating a lower-resolution "view".

The final discriminator is a weighted sum over all outputs,

$$D(\bar{x}) = \sum_s \alpha_s T^{s \times s} \left(D^s(T^{s^{-1} \times s^{-1}}(\bar{x})) \right),$$

Step 2: Patchwise Probability Calculation per Sub-Discriminator

Each sub-discriminator D_s is a **convolutional network** (e.g., a series of 4×4 strided conv layers).

- **Input:** The resized spectrogram from T_s of size [1, Freq/ s , Time/ s].
- **Process:** The convolutional network processes this input. The key design is that the **final layer is a convolution (not a linear layer)**.
- **Output:** The output of D_s is not a single number. It is a **2D feature map** (e.g., of size [1, 4, 8]). Each value in this [4, 8] grid corresponds to the **logit** for a specific "patch" or receptive field in the *resized* input.
- **Patchwise Reality:** This means D_s produces a **dense map of "realness" scores**. It can identify that, e.g., the patch covering timeframes 10-20 and frequency bins 40-50 looks fake, while the rest looks real. This provides a incredibly rich and precise gradient signal for the generator.

Step 3: Combining Scales via Weighted Sum

The outputs of all 5 sub-discriminators are not independent; they are combined into a final, unified "realness" map.

1. **Upsample Discriminator Outputs:** The output of each D_s (e.g., the [4, 8] grid) To combine them, we must **upsample each output back to the original resolution** of the spectrogram being judged.

- We apply the inverse transformation: T_s^{-1} .
- This results in 5 matrices ($0_s=1, 0_s=1.2, \dots, 0_s=2.07$), all at the original resolution (e.g., [80, 384] for Y_{fake}), representing the realness score for every potential patch.

2. **Learnable Weighted Sum:** The final discriminator output is a weighted sum of these upsampled maps:

- α_s are **learnable parameters**.
- The network learns **how much to trust the judgment of each scale**. For example, it might learn that the finest scale ($s=1$) is most important for judging harmonic sharpness, while a mid-scale ($s=1.44$) is best for judging rhythmic consistency.

Training Objective: Least-Squares GAN (LSGAN) Loss

Discriminator's Goal: Maximize $L^{L_s}(G, D)$

The discriminator D aims to maximize its ability to distinguish real from fake. This is done by minimizing the following loss function:

$$L^{L_s}(G, D) = E_{\tilde{x}}[pdata(\tilde{x})[(D(\tilde{x}) - J)^2] + E_{\tilde{x}, r}[pdata(\tilde{x}, r)[(D(G(\tilde{x}, r)) - 0)^2]]$$

Term 1: Real Loss (L_{real})

$$E_{\tilde{x}}[pdata(\tilde{x})[(D(\tilde{x}) - J)^2]$$

- **Calculation:** For a batch of real spectrograms, the discriminator's output $D(\text{real})$ is a matrix of patch-wise predictions. We calculate the Mean Squared Error (MSE) between this matrix and a matrix of all ones (J).
- **Interpretation:** This loss is minimized when $D(\text{real})$ outputs values close to 1 for every patch. It teaches the discriminator to assign the "real" label to ground-truth data.

Term 2: Fake Loss (L_{fake})

$$E_{\tilde{x}, r}[(D(G(\tilde{x}, r)) - 0)^2]$$

- **Calculation:** For the same batch of real spectrograms, we pass them through the generator G with a random rate r to get fake spectrograms $G(\text{real}, r)$. The discriminator's output $D(\text{fake})$ is compared to a matrix of all zeros (0) using MSE.
- **Interpretation:** This loss is minimized when $D(\text{fake})$ outputs values close to 0 for every patch. It teaches the discriminator to assign the "fake" label to generated data.

$$\text{Discriminator's Total Loss: } L^D = L_{real} + L_{fake}$$

The discriminator's weights are updated to minimize L^D .

2. Generator's Goal: Minimize $L^{ls}(G, D)$

The generator G aims to minimize the discriminator's ability to catch fakes. It does this by minimizing its own specific loss function:

$$L^G = E_{\bar{x}, r}[(D(G(\bar{x}, r)) - J)^2]$$

- **Calculation:** The generator's loss is calculated on the same batch of fake spectrograms. However, here we calculate the MSE between $D(\text{fake})$ and a matrix of all ones (J).
- **Interpretation:** This loss is minimized when $D(\text{fake})$ is fooled into outputting values close to 1. It teaches the generator to produce outputs that are indistinguishable from real data, according to the discriminator.

Generator's Adversarial Loss: $L^G_{\text{adv}} = E[(D(G(\bar{x}, r)) - J)^2]$

The generator's weights are updated to minimize L^G_{adv} .

3. The Full Objective & Weight Update Process

The complete training goal is a minimax game with an additional constraint:

$$\min^G \max^D [L^{ls}(G, D) + \lambda L^R(G)], \lambda = 0.1$$

Where $L^R(G)$ is the Cycle-Consistency Reconstruction Loss (e.g., L1 loss between original and reconstructed spectrogram).

Reconstruction Loss

- Cycle consistency:
We first time-scale $\mathbf{x} \rightarrow \mathbf{y} = G(\mathbf{x}, r)$, then scale it back $\mathbf{y} \rightarrow \mathbf{x}' = G(\mathbf{y}, 1/r)$
- Goal:
Ensure that reconstructed \mathbf{x}' is close to the original \mathbf{x}
- Why?
Prevents mode collapse (e.g., G generating repetitive or blurry outputs)

The second loss function is associated with the decoder. We would like minimize the L_1 distance between $\bar{\mathbf{x}}' = G(\bar{\mathbf{y}}, r^{-1})$ and $\bar{\mathbf{x}}$ so as to prevent convergence to trivial solutions:

$$\mathcal{L}_R(G) = \|G(G(\bar{\mathbf{x}}, r), r^{-1}) - \bar{\mathbf{x}}\|_1 , \quad (2)$$

Experiments and Results

1. Datasets & Preprocessing

- **Primary Dataset: LJSpeech**
 - 13,100 utterances (24 hrs), single female speaker, 22.05 kHz, 16-bit PCM
- **Generalization Test: VCTK Subset**
 - 28 speakers (14 M, 14 F) → Evaluates unseen voice adaptation
- **Feature Extraction (Mel-spectrograms):**
 - `n_fft = 1024, hop_size = 256, n_mels = 80`
 - `fmin = 0, fmax = 8000` (Covers speech-relevant frequencies)

4. Key Results

- **Successful Convergence (Epoch 500):**
 - `G_loss = 0.2486`
 - `D_loss_real = 0.2496`
 - `D_loss_fake = 0.2533`
 - `Rec_loss = 0.1111`
- **Qualitative Outputs:**
 - Temporal Scaling → Generator adjusts duration while preserving harmonics (Fig. 4)
 - Cycle Consistency → Low reconstruction error (Fig. 5), content identity maintained
 - Loss Curves → Stable adversarial training (Fig. 6)
 - Generalization → Effective scaling on unseen VCTK speakers


2. Experimental Setup

- **Training:** Random segments of 256 frames
- **Inference:** Full spectrograms
- **Architecture:**
 - Generator: U-Net with 6 residual blocks (bottleneck)
 - Discriminator: 5 sub-discriminators (multi-scale)
- **Hyperparameters:**
 - Batch Size = 24 | Learning Rate = 5e-5 (Adam, $\beta_1=0.5, \beta_2=0.999$)
 - Epochs = 500 | λ (Reconstruction Weight) = 0.1
- **Curriculum Learning:**
 - Scaling factor r gradually expanded from [1.0, 1.0] → [0.3, 1.8]
- **Stability Measures:**
 - Checkpoint resuming
 - Inverse steps added post-epoch 200 (cycle consistency)


4. Inference Phase

- At test time, the generated spectrogram \bar{y} is converted back into an audio waveform using a neural vocoder (HiFi-GAN)


mel



g_pred




origin mel



difference

reconstruct



FastSpeech: Ideas for Content-Aware TSM

Limitations of Uniform Time-Scaling

ScalerGAN performs **content-agnostic scaling**, applying the same temporal transformation globally. This can introduce artifacts:

- **Over-compression** of long phonemes (e.g., vowels)
- **Over-extension** of short phonemes (e.g., plosives: /p/, /t/, /k/)
- **Loss of prosody**: Natural rhythm, stress, and pauses are not preserved
- **Reduced naturalness** despite perceptual realism in spectrograms

Proposed Solution: Integrate FastSpeech Components

To enable **phoneme-aware duration control**, we propose hybridizing ScalerGAN with two key modules from **FastSpeech**:

1. Duration Predictor (DP)

- A lightweight convolutional network
- **Predicts natural phoneme durations** (in frames) from input features
- Trained using alignments from a pre-trained teacher model (e.g., Tacotron 2)

Example:

- Phonemes: [p,o,t]
- Durations: [2,5,3] → Expanded: [p,p,o,o,o,o,o,t,t,t]

2. Length Regulator (LR)

- **Expands/compresses phoneme-level features** based on predicted durations
- Uses a trainable duration scale factor
- (α) for global speed control

Thank You